

Assignment 20

This Assignment contains some code blocks took from original ipython notebooks (4.ML_models.ipynb and other notebooks).

I am using following files that are created from other ipython notebooks to run the models.

- train.csv
- nlp_features_train.csv
- df_fe_without_preprocessing_train.csv
- train.db

```
In [2]: # Required Imports
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
import datetime as dt
import numpy as np

from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
```

```
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve
```

/usr/local/lib/python3.6/dist-packages/sklearn/externals/six.py:31: DeprecationWarning: The module is deprecated in version 0.21 and will be removed in version 0.23 since we've dropped support for Python 2.7. Please rely on the official version of six (<https://pypi.org/project/six/>).

"(<https://pypi.org/project/six/>).", DeprecationWarning)

Task 1

Train Logistic Regression and Linear-SVM with TfIdf vectorization of questions data.

Taking only first 50,000 rows to train and test the data due to computation limitations.

```
In [3]: main_data = pd.read_csv('train.csv', nrows=50000)
main_data.head()
```

```
Out[3]:
```

	id	qid1	qid2	question1	question2	is_duplicate
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when 23^{24} $\pmod{100}$	0
4	4	9	10	Which one dissolve in water quickly sugar, salt...	Which fish would survive in salt water?	0

```
In [4]: nlp_feats = pd.read_csv('nlp_features_train.csv', nrows=50000, encoding='latin-1')
nlp_feats.head(3)
```

```
Out[4]:
```

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_i
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	0.916659	0.785
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	0.699993	0.466
2	2	5	6	how can i increase the speed of my internet co...	how can internet speed be increased by hacking...	0	0.399992	0.333328	0.399992	0.249997	0.399996	0.285

```
In [5]: feat_engg = pd.read_csv('df_fe_without_preprocessing_train.csv', nrows=50000, encoding='lat
feat_engg.head(3)
```

Out[5]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_w
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0	1	1	73	59	14	

Splitting data into train and test data before applying Tfidf vectorizer to avoid data leakage

```
In [6]: questions = main_data[['question1', 'question2']]
labels = pd.DataFrame(main_data['is_duplicate'])
questions.head()
```

```
Out[6]:
```

	question1	question2
0	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...
1	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...
2	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...
3	Why am I mentally very lonely? How can I solve...	Find the remainder when 23^{24} i...
4	Which one dissolve in water quikly sugar, salt...	Which fish would survive in salt water?

```
In [0]: q_train, q_test, labels_train, labels_test = train_test_split(questions, labels, test_size=
```

```
In [0]: train_ind = q_train.index
test_ind = q_test.index
```

```
In [9]: from sklearn.feature_extraction.text import TfidfVectorizer
quest_to_vectorize = list(q_train['question1']) + list(q_train['question2'])
tfidf = TfidfVectorizer(stop_words='english', max_features=5000)
tfidf.fit(quest_to_vectorize)
```

```
Out[9]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.float64'>, encoding='utf-8',
input='content', lowercase=True, max_df=1.0, max_features=5000,
min_df=1, ngram_range=(1, 1), norm='l2', preprocessor=None,
smooth_idf=True, stop_words='english', strip_accents=None,
sublinear_tf=False, token_pattern='(?u)\\b\\w+\\b',
tokenizer=None, use_idf=True, vocabulary=None)
```

```
In [10]: q1_train = tfidf.transform(q_train['question1'].values)
q2_train = tfidf.transform(q_train['question2'].values)
q1_test = tfidf.transform(q_test['question1'].values)
q2_test = tfidf.transform(q_test['question2'].values)

print(q1_train.shape)
print(q2_train.shape)
print(q1_test.shape)
print(q2_test.shape)

(40000, 5000)
(40000, 5000)
(10000, 5000)
(10000, 5000)
```

Splitting all DataFrames according to train_ind and test_ind obtained from train_test_split and combine them to form X_train and X_test

```
In [11]: feat_engg.drop(columns=main_data.columns, inplace=True)
nlp_feats.drop(columns=main_data.columns, inplace=True)
print(feat_engg.shape)
print(nlp_feats.shape)

(50000, 11)
(50000, 15)
```



```
In [12]: feat_engg_train = feat_engg.loc[train_ind, :]  
feat_engg_test = feat_engg.loc[test_ind, :]  
nlp_feats_train = nlp_feats.loc[train_ind, :]  
nlp_feats_test = nlp_feats.loc[test_ind, :]
```

```
print(feat_engg_train.shape)  
print(feat_engg_test.shape)  
print(nlp_feats_train.shape)  
print(nlp_feats_test.shape)
```

```
(40000, 11)  
(10000, 11)  
(40000, 15)  
(10000, 15)
```

```
In [13]: X_train = hstack((feat_engg_train, nlp_feats_train, q1_train, q2_train)).tocsr()  
X_test = hstack((feat_engg_test, nlp_feats_test, q1_test, q2_test)).tocsr()  
y_train = labels_train  
y_test = labels_test
```

```
print(X_train.shape, y_train.shape)  
print(X_test.shape, y_test.shape)
```

```
(40000, 10026) (40000, 1)  
(10000, 10026) (10000, 1)
```

As our data is ready we perform Logistic Regression and Linear-SVM on the data. For this I will take the code from one of the original ipython notebooks

```

In [0]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in that column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1)  axis=0 corresonds to columns and axis=1 corresponds to rows in two d
    # C.sum(axix =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B =(C/C.sum(axis=0))
    #divid each element of the confusion matrix with the sum of elements in that row
    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0)  axis=0 corresonds to columns and axis=1 corresponds to rows in two d
    # C.sum(axix =0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]
    plt.figure(figsize=(20,4))

    labels = [1,2]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)

```

```
sns.heatmap(C, annot=True, cmap=cmap, fmt=".0f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")


plt.show()
```

Training a Random model to know worst-log-loss that we can get.

```
In [15]: test_len = len(y_test)
train_len = len(y_train)
predicted_y = np.zeros((test_len,2))
for i in range(test_len):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test, predicted_y, eps=1e-15))

predicted_y =np.argmax(predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y)
```

Log loss on Test Data using Random Model 0.8974195927624847



Logistic Regression with hyper-parameter tuning

```

In [16]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, label

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))


```

```
plot_confusion_matrix(y_test, predicted_y)
```

For values of alpha = 1e-05 The log loss is: 0.427303052937189
For values of alpha = 0.0001 The log loss is: 0.4146490324907647
For values of alpha = 0.001 The log loss is: 0.42354989460930553
For values of alpha = 0.01 The log loss is: 0.4376423906863354
For values of alpha = 0.1 The log loss is: 0.45904348314400834
For values of alpha = 1 The log loss is: 0.4906774962637248
For values of alpha = 10 The log loss is: 0.5459613692305979



For values of best alpha = 0.0001 The train log loss is: 0.4127349578718349
For values of best alpha = 0.0001 The test log loss is: 0.4146490324907647
Total number of data points : 10000



The loss values for Tfidf vectorized data is much less than Tfidf Word2Vec loss values for Logistic Regression. Which means this model did lot better than the model in original notebook

Linear-SVM with hyper-parameter tuning

```
In [17]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, label

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()


best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
```




```
plot_confusion_matrix(y_test, predicted_y)
```

For values of alpha = 1e-05 The log loss is: 0.42085570185631943
For values of alpha = 0.0001 The log loss is: 0.44496904081309385
For values of alpha = 0.001 The log loss is: 0.4756987604337793
For values of alpha = 0.01 The log loss is: 0.5171187047471906
For values of alpha = 0.1 The log loss is: 0.47957363952464294
For values of alpha = 1 The log loss is: 0.573956591705653
For values of alpha = 10 The log loss is: 0.5969699647626083



For values of best alpha = 1e-05 The train log loss is: 0.4154266037600815
For values of best alpha = 1e-05 The test log loss is: 0.42085570185631943
Total number of data points : 10000



Here also we can see that our model improved with Tfidf vectorization in Linear-SVM. and also we can observe in both Logistic Regression and Linear-SVM that they are not overfitted as the train loss and test loss are almost equal

Task 2

Do hyper-parameter tuning on XGBoost model with Tfidf Word2Vec data

Vectorizing the data to Tfidf Word2Vec again as the models in original notebooks have data leakage

```
In [0]: import spacy
        from tqdm import tqdm
```

```
In [0]: tfidf = TfidfVectorizer(stop_words='english')
        tfidf.fit_transform(quest_to_vectorize) # these questions are only taken from train data

        # dict key:word and value:tf-idf score
        word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))
```

```
In [0]: # en_vectors_web_lg, which includes over 1 million unique vectors.
nlp = spacy.load('en_core_web_sm')
```

```
def tfidf_w2v_transform(data):
    vecs = []
    # https://github.com/noamraph/tqdm
    # tqdm is used to print the progress bar
    for qu in tqdm(list(data)):
        doc = nlp(qu)
        # 384 is the number of dimensions of vectors
        mean_vec = np.zeros([len(doc), len(doc[0].vector)])
        for word in doc:
            # word2vec
            vec = word.vector
            # fetch df score
            try:
                idf = word2tfidf[str(word)]
            except:
                idf = 0
            # compute final vec
            mean_vec += vec * idf
        mean_vec = mean_vec.mean(axis=0)
        vecs.append(mean_vec)
    return list(vecs)
```

```
In [21]: q1_w2v_train = tfidf_w2v_transform(q_train['question1'].values)
q2_w2v_train = tfidf_w2v_transform(q_train['question2'].values)
q1_w2v_test = tfidf_w2v_transform(q_test['question1'].values)
q2_w2v_test = tfidf_w2v_transform(q_test['question2'].values)
```

```
100%|██████████| 40000/40000 [07:01<00:00, 94.86it/s]
100%|██████████| 40000/40000 [07:01<00:00, 95.01it/s]
100%|██████████| 10000/10000 [01:44<00:00, 95.81it/s]
100%|██████████| 10000/10000 [01:44<00:00, 95.29it/s]
```

```
In [22]: X_w2v_train = hstack((feat_engg_train, nlp_feats_train, q1_w2v_train, q2_w2v_train)).tocsr()  
X_w2v_test = hstack((feat_engg_test, nlp_feats_test, q1_w2v_test, q2_w2v_test)).tocsr()  
  
print(X_w2v_train.shape, y_train.shape)  
print(X_w2v_test.shape, y_test.shape)
```

```
(40000, 218) (40000, 1)  
(10000, 218) (10000, 1)
```

The scapy nlp function gives a vector of 96 dimention for a word. So the dimentions are not as same as the dimentions in original notebook

```
In [23]: len(nlp('Hello').vector)
```

```
Out[23]: 96
```

XGBoost with hyper-parameter tuning

```
In [24]: import xgboost as xgb

d_train = xgb.DMatrix(X_w2v_train, label=y_train)
d_test = xgb.DMatrix(X_w2v_test, label=y_test)

watchlist = [(d_train, 'train'), (d_test, 'valid')]

train_result = []
test_result = []

etas = [0.02, 0.05, 0.1, 0.2, 0.3]
depths = [4, 6, 8] # Taking limited range for XGBoost as it only requires underfitted base

for eta in etas:
    train_sub_res = []
    test_sub_res = []
    for dep in depths:
        params = {}
        params['objective'] = 'binary:logistic'
        params['eval_metric'] = 'logloss'
        params['eta'] = eta
        params['max_depth'] = dep

    bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_eval=False)

    xgdmatrix = xgb.DMatrix(X_w2v_train, y_train)

    predict_y = bst.predict(d_train)
    train_loss = log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15)
    train_sub_res.append(train_loss)

    predict_y = bst.predict(d_test)
    test_loss = log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15)
    test_sub_res.append(test_loss)
```

```
print(f"eta = {eta} and max_depth = {dep}: train loss = {train_loss} and test loss = {t
train_result.append(train_sub_res)
test_result.append(test_sub_res)

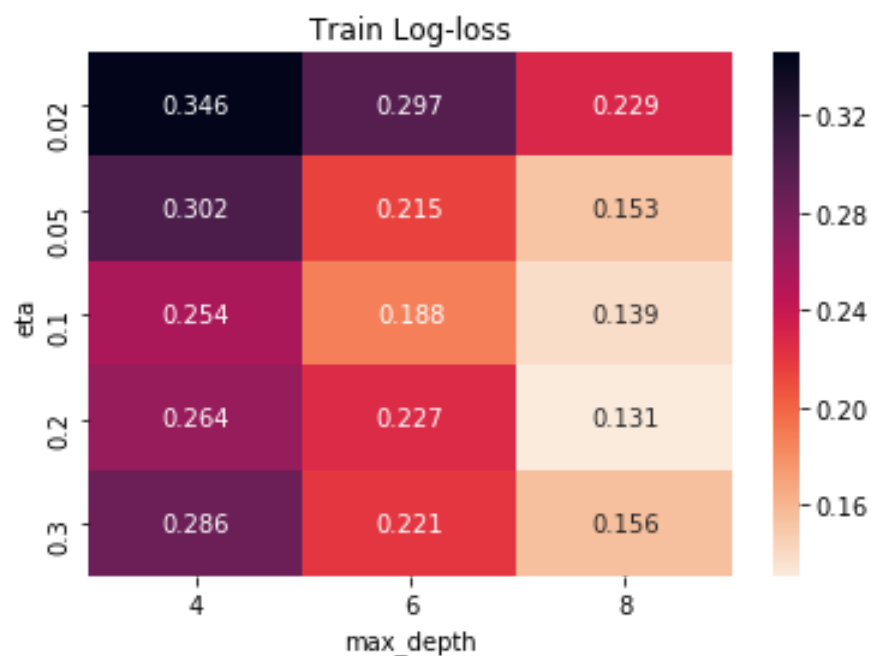
print(train_result)
print(test_result)
```

```
eta = 0.02 and max_depth = 4: train loss = 0.3456527848904836 and test loss = 0.357774577
8926357
eta = 0.02 and max_depth = 6: train loss = 0.29704844136157044 and test loss = 0.34710550
052327455
eta = 0.02 and max_depth = 8: train loss = 0.22904690350333987 and test loss = 0.34219318
75396171
eta = 0.05 and max_depth = 4: train loss = 0.3022930904685154 and test loss = 0.344007251
9325324
eta = 0.05 and max_depth = 6: train loss = 0.21471708702347642 and test loss = 0.34037295
57209994
eta = 0.05 and max_depth = 8: train loss = 0.15272361931330034 and test loss = 0.33879815
431861715
eta = 0.1 and max_depth = 4: train loss = 0.2542708137487427 and test loss = 0.3383732428
0897146
eta = 0.1 and max_depth = 6: train loss = 0.18800298595766268 and test loss = 0.339838671
58069025
eta = 0.1 and max_depth = 8: train loss = 0.13884129914616386 and test loss = 0.342545721
5739836
eta = 0.2 and max_depth = 4: train loss = 0.2642600291833135 and test loss = 0.3438030465
3407545
eta = 0.2 and max_depth = 6: train loss = 0.22678442535402377 and test loss = 0.348596406
9886304
eta = 0.2 and max_depth = 8: train loss = 0.13074296321007114 and test loss = 0.351473090
8657395
eta = 0.3 and max_depth = 4: train loss = 0.28628159365964295 and test loss = 0.346135359
1497171
eta = 0.3 and max_depth = 6: train loss = 0.22113488655455116 and test loss = 0.353240745
8626176
```

```
eta = 0.3 and max_depth = 8: train loss = 0.15560197392924002 and test loss = 0.35618309074652327  
[[0.3456527848904836, 0.29704844136157044, 0.22904690350333987], [0.3022930904685154, 0.21471708702347642, 0.15272361931330034], [0.2542708137487427, 0.18800298595766268, 0.13884129914616386], [0.2642600291833135, 0.22678442535402377, 0.13074296321007114], [0.28628159365964295, 0.22113488655455116, 0.15560197392924002]]  
[[0.3577745778926357, 0.34710550052327455, 0.3421931875396171], [0.3440072519325324, 0.3403729557209994, 0.33879815431861715], [0.33837324280897146, 0.33983867158069025, 0.3425457215739836], [0.34380304653407545, 0.3485964069886304, 0.3514730908657395], [0.3461353591497171, 0.3532407458626176, 0.35618309074652327]]
```

```
In [30]: sns.heatmap(train_result, annot=True, fmt='.3f', xticklabels=depths, yticklabels=etas, cmap=
plt.title('Train Log-loss')
plt.xlabel('max_depth')
plt.ylabel('eta')
plt.show()

sns.heatmap(test_result, annot=True, fmt='.3f', xticklabels=depths, yticklabels=etas, cmap=
plt.title('Test Log-loss')
plt.xlabel('max_depth')
plt.ylabel('eta')
plt.show()
```





All the models are overfitting except with parameters max_depth=4 and eta=0.02, 0.05. As we saw performance of eta=0.02 and depth=4 in original notebook, I am going to train and show results for eta=0.05 and depth=4

```
In [31]: params = {}
params['objective'] = 'binary:logistic'
params['eval_metric'] = 'logloss'
params['eta'] = 0.05
params['max_depth'] = 4

# dtrain and watchlist is defined in above code blocks
bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_eval=10)

predict_y = bst.predict(d_test)
print("The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```
[0]    train-logloss:0.672729  valid-logloss:0.672858
```

Multiple eval metrics have been passed: 'valid-logloss' will be used for early stopping.


Will train until valid-logloss hasn't improved in 20 rounds.

```
[10]    train-logloss:0.53703  valid-logloss:0.537403
[20]    train-logloss:0.469865  valid-logloss:0.470636
[30]    train-logloss:0.432191  valid-logloss:0.433207
[40]    train-logloss:0.409219  valid-logloss:0.410467
[50]    train-logloss:0.394428  valid-logloss:0.396104
[60]    train-logloss:0.384739  valid-logloss:0.38692
[70]    train-logloss:0.377629  valid-logloss:0.380617
[80]    train-logloss:0.372184  valid-logloss:0.376053
[90]    train-logloss:0.368042  valid-logloss:0.372772
[100]   train-logloss:0.364294  valid-logloss:0.369606
[110]   train-logloss:0.360393  valid-logloss:0.366693
[120]   train-logloss:0.356797  valid-logloss:0.364178
[130]   train-logloss:0.353702  valid-logloss:0.362361
[140]   train-logloss:0.350662  valid-logloss:0.360611
[150]   train-logloss:0.347745  valid-logloss:0.358858
[160]   train-logloss:0.345049  valid-logloss:0.357359
[170]   train-logloss:0.342588  valid-logloss:0.355993
[180]   train-logloss:0.340279  valid-logloss:0.355191
[190]   train-logloss:0.337944  valid-logloss:0.354066
```

```
[200] train-logloss:0.335594 valid-logloss:0.353083
[210] train-logloss:0.333389 valid-logloss:0.352165
[220] train-logloss:0.331383 valid-logloss:0.351546
[230] train-logloss:0.329137 valid-logloss:0.350705
[240] train-logloss:0.327256 valid-logloss:0.349933
[250] train-logloss:0.325649 valid-logloss:0.349574
[260] train-logloss:0.324141 valid-logloss:0.349191
[270] train-logloss:0.322172 valid-logloss:0.348568
[280] train-logloss:0.320592 valid-logloss:0.348147
[290] train-logloss:0.318669 valid-logloss:0.347409
[300] train-logloss:0.317183 valid-logloss:0.347089
[310] train-logloss:0.31548 valid-logloss:0.346531
[320] train-logloss:0.313975 valid-logloss:0.345969
[330] train-logloss:0.312515 valid-logloss:0.345834
[340] train-logloss:0.310704 valid-logloss:0.345286
[350] train-logloss:0.309287 valid-logloss:0.345013
[360] train-logloss:0.307923 valid-logloss:0.344744
[370] train-logloss:0.306493 valid-logloss:0.344647
[380] train-logloss:0.305216 valid-logloss:0.344303
[390] train-logloss:0.303763 valid-logloss:0.344176
[399] train-logloss:0.302288 valid-logloss:0.344008
The test log loss is: 0.3440072519325324
```

```
In [32]: predicted_y = np.array(predict_y > 0.5, dtype=int)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)
```

Total number of data points : 10000



Conclusion

- Linear Regression and Linear-SVM models did very good with Tfidf vectorized data. And XGBoost model did slightly better with hyper-parameter tuning, as the selected hyper-parameters in original notebook are near to optimal ones.
- XGBoost has lot of overfitting even when our max_depth is low. Until now the XGBoost model gave good performance compared to other models we tried.
- The reason for good performance of Linear models with Tfidf Vectorization is due to high dimensionality.

In [0]:

